



FOODIE

DESIGN NOTES

Solving for resolution & accuracy in B2B conversational marketplaces.

INTERVIEW SEGMENT

Software Engineering Intro


DURATION

3-Minute Technical Demo

FOCUS

Hard Problems & Workflow

Resolution: one tool, four outcomes.



HOW IT WORKS

The voice agent calls one tool: `resolve_item`. The logic runs in Python — not in the LLM.

- Deterministic
- Unit-testable
- Catalog Isolation

FUTURE EXTENSIONS

Sponsored brands

Suppliers could pay to surface in the recommendation slot. Commercial priority vs habit.

Collaborative filtering

Recommend based on similar customers when no history exists.

So the LLM can't hallucinate a SKU we don't carry.

Confirm what we heard. Handle what we didn't.

CONFIRMATION STRATEGY



Per-item echo

"Tomato sauce, got it — how many?" Catches mishears immediately. Zero overhead as we need quantity anyway.



Full recap before save

Every item + brand + quantity. One explicit "yes" required before the `save_order` tool fires.

FAILURE RECOVERY



Item not in catalog

Agent is instructed to ask for clarification rather than guessing a near-match from the catalog.



Vague quantity

Clarify immediately. Never assume "a couple" = 2. Validates against integer constraints.

Every uncertainty is either resolved or made visible — never silently assumed.

Compliance, built in.

Australian telemarketing rules, enforced before we dial.



Do Not Call — Absolute

Customer-level flag checked before every dial. Cannot be bypassed, ensuring strict adherence to user preferences even in demos.



Calling Hours

Timezone-aware windows (Australia/Sydney) per the 2017 Industry Standard. Enforced at the dialer level, not the prompt level.



Opt-out — First-class

`flag_do_not_call` is a core tool. Real-time flagging mid-call immediately updates the dashboard UI to prevent future dials.

Not legal advice — design constraints, built in from day one.